

Pessimistic Query Optimization: Tighter Upper Bounds for Intermediate Join Cardinalities

Walter Cai Magdalena Balazinska Dan Suciu

University of Washington

[walter,magda,suciu]@cs.washington.edu

April 23, 2019

Contributions

- ▶ Technique for tightening theoretically guaranteed join cardinality upper bounds.

Contributions

- ▶ Technique for tightening theoretically guaranteed join cardinality upper bounds.
- ▶ Method for enumerating practical subset of bounding formulas.

Contributions

- ▶ Technique for tightening theoretically guaranteed join cardinality upper bounds.
- ▶ Method for enumerating practical subset of bounding formulas.
- ▶ Partition budgeting strategy to control the space complexity of our sketches, and the time complexity of our bound calculation.

Contributions

- ▶ Technique for tightening theoretically guaranteed join cardinality upper bounds.
- ▶ Method for enumerating practical subset of bounding formulas.
- ▶ Partition budgeting strategy to control the space complexity of our sketches, and the time complexity of our bound calculation.
- ▶ Demonstrate practicality on challenging real world benchmark.

- 1 Query Optimization
- 2 Motivating Example
- 3 Prior Work: Cardinality Bounds
- 4 Tightened Cardinality Bounds
- 5 Optimizations
- 6 Evaluation
 - Bound Tightening
 - Runtime Improvement
- 7 Conclusion and Future Directions

Query Optimization

- ▶ Accepts queries.
- ▶ Picks “best” physical plan.
 - ▶ Could be millions of correct physical plans!
 - ▶ Conceptually, a tree with leaves as base relations.


Query Optimization

- ▶ Accepts queries.
- ▶ Picks “best” physical plan.
 - ▶ Could be millions of correct physical plans!
 - ▶ Conceptually, a tree with leaves as base relations.

cast


Query Optimization

- ▶ Accepts queries.
- ▶ Picks “best” physical plan.
 - ▶ Could be millions of correct physical plans!
 - ▶ Conceptually, a tree with leaves as base relations.



Query Optimization

- ▶ Accepts queries.
- ▶ Picks “best” physical plan.
 - ▶ Could be millions of correct physical plans!
 - ▶ Conceptually, a tree with leaves as base relations.



Query Optimization

- ▶ Accepts queries.
- ▶ Picks “best” physical plan.
 - ▶ Could be millions of correct physical plans!
 - ▶ Conceptually, a tree with leaves as base relations.




Figure: Join tree illustrations.

Query Optimization

- ▶ Cost-Based.
 - ▶ Large parameterized summation.
 - ▶ Sum over cost of each physical operator.




Figure: Join tree illustrations.

Query Optimization

- ▶ Join Algorithms are generally binary so the DBMS will generate intermediate relations.
- ▶ **Cardinality Estimation:** how large will these intermediate relations be?




Figure: Join tree illustrations.

Cardinality Estimation Error

- ▶ Systems rely on strong assumptions about the underlying data.

Cardinality Estimation Error

- ▶ Systems rely on strong assumptions about the underlying data.
- ▶ Assume independence of attribute value distributions across columns.

Cardinality Estimation Error

- ▶ Systems rely on strong assumptions about the underlying data.
- ▶ Assume independence of attribute value distributions across columns.
- ▶ Leads to underestimation.
 - ▶ Real world data is correlated.
 - ▶ Underestimation is risky: leads to massive blow-up from poor join orderings/algorithm choice.

- 1 Query Optimization
- 2 Motivating Example**
- 3 Prior Work: Cardinality Bounds
- 4 Tightened Cardinality Bounds
- 5 Optimizations
- 6 Evaluation
 - Bound Tightening
 - Runtime Improvement
- 7 Conclusion and Future Directions

Join Order Benchmark (JOB)

- ▶ Built on the IMDb dataset.

Join Order Benchmark (JOB)

- ▶ Built on the IMDb dataset.
 - ▶ 113 queries.

Join Order Benchmark (JOB)

- ▶ Built on the IMDb dataset.
 - ▶ 113 queries.
 - ▶ 33 unique topologies.

Join Order Benchmark (JOB)

- ▶ Built on the IMDb dataset.
 - ▶ 113 queries.
 - ▶ 33 unique topologies.
 - ▶ Skew!

Join Order Benchmark (JOB)

- ▶ Built on the IMDb dataset.
 - ▶ 113 queries.
 - ▶ 33 unique topologies.
 - ▶ Skew!
 - ▶ Correlation!

Join Order Benchmark (JOB)

- ▶ Built on the IMDb dataset.
 - ▶ 113 queries.
 - ▶ 33 unique topologies.
 - ▶ Skew!
 - ▶ Correlation!
 - ▶ Complex selection predicates!


JOB Example Query

```

SELECT
    *
FROM
    aka,
    cast,
    company_name,
    movie_companies,
    name,
    role,
    title
WHERE
    company_name.country = 'usa' AND
    role.type = 'writer' AND
    aka.person_id = name.id AND
    cast.person_id = name.id AND
    aka.person_id = cast.person_id AND
    cast.movie_id = title.id AND
    movie_companies.movie = title.id_id AND
    cast.movie_id = movie_companies.movie_id AND
    movie_companies.company_id = company_name.id AND
    cast.role_id = role.id;

```


JOB Example Query



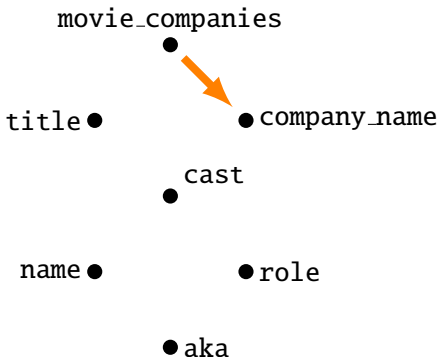
```

SELECT
    *
FROM
    aka,
    cast,
    company_name,
    movie_companies,
    name,
    role,
    title
WHERE
    company_name.country = 'usa' AND
    role.type = 'writer' AND
    aka.person_id = name.id AND
    cast.person_id = name.id AND
    aka.person_id = cast.person_id AND
    cast.movie_id = title.id AND
    movie_companies.movie = title.id AND
    cast.movie_id = movie_companies.movie_id AND
    movie_companies.company_id = company_name.id AND
    cast.role_id = role.id;

```

Figure: Join Graph.

JOB Example Query



```

SELECT
  *
FROM
  aka,
  cast,
  company_name,
  movie_companies,
  name,
  role,
  title
WHERE
  company_name.country = 'usa' AND
  role.type = 'writer' AND
  aka.person_id = name.id AND
  cast.person_id = name.id AND
  aka.person_id = cast.person_id AND
  cast.movie_id = title.id AND
  movie_companies.movie = title.id_id AND
  cast.movie_id = movie_companies.movie_id AND
  movie_companies.company_id = company_name.id AND
  cast.role_id = role.id;

```

Figure: Join Graph.

JOB Example Query

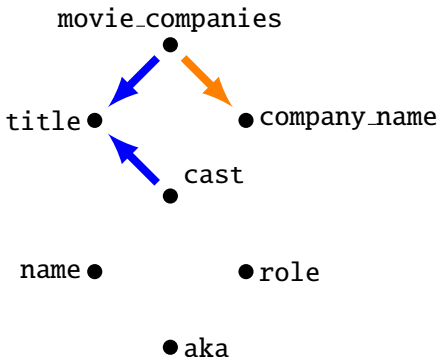


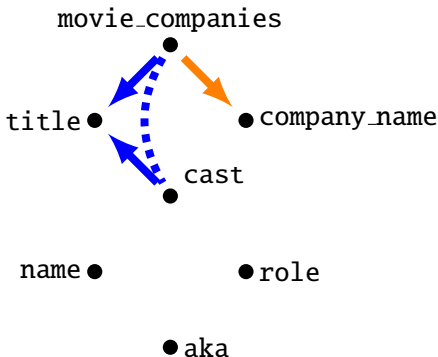
Figure: Join Graph.

```

SELECT
  *
FROM
  aka,
  cast,
  company_name,
  movie_companies,
  name,
  role,
  title
WHERE
  company_name.country = 'usa' AND
  role.type = 'writer' AND
  aka.person_id = name.id AND
  cast.person_id = name.id AND
  aka.person_id = cast.person_id AND
  cast.movie_id = title.id AND
  movie_companies.movie = title.id_id AND
  cast.movie_id = movie_companies.movie_id AND
  movie_companies.company_id = company_name.id AND
  cast.role_id = role.id;

```

JOB Example Query



```

SELECT
  *
FROM
  aka,
  cast,
  company_name,
  movie_companies,
  name,
  role,
  title
WHERE
  company_name.country = 'usa' AND
  role.type = 'writer' AND
  aka.person_id = name.id AND
  cast.person_id = name.id AND
  aka.person_id = cast.person_id AND
  cast.movie_id = title.id AND
  movie_companies.movie = title.id_id AND
  cast.movie_id = movie_companies.movie_id AND
  movie_companies.company_id = company_name.id AND
  cast.role_id = role.id;

```

Figure: Join Graph.

JOB Example Query

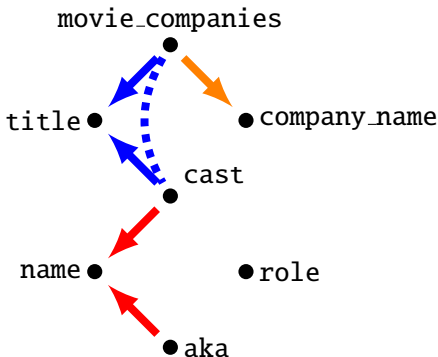


Figure: Join Graph.

```

SELECT
  *
FROM
  aka,
  cast,
  company_name,
  movie_companies,
  name,
  role,
  title
WHERE
  company_name.country = 'usa' AND
  role.type = 'writer' AND
  aka.person_id = name.id AND
  cast.person_id = name.id AND
  aka.person_id = cast.person_id AND
  cast.movie_id = title.id AND
  movie_companies.movie = title.id_id AND
  cast.movie_id = movie_companies.movie_id AND
  movie_companies.company_id = company_name.id AND
  cast.role_id = role.id;

```

JOB Example Query




Figure: Join Graph.

```

SELECT
  *
FROM
  aka,
  cast,
  company_name,
  movie_companies,
  name,
  role,
  title
WHERE
  company_name.country = 'usa' AND
  role.type = 'writer' AND
  aka.person_id = name.id AND
  cast.person_id = name.id AND
  aka.person_id = cast.person_id AND
  cast.movie_id = title.id AND
  movie_companies.movie = title.id_id AND
  cast.movie_id = movie_companies.movie_id AND
  movie_companies.company_id = company_name.id AND
  cast.role_id = role.id;

```

JOB Example Query

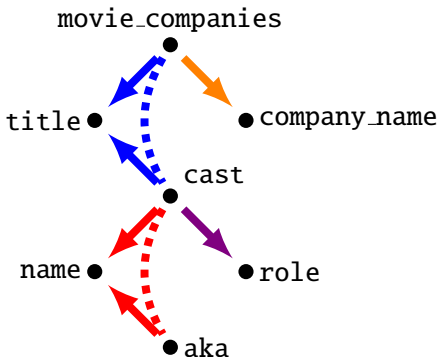



Figure: Join Graph.

```

SELECT
  *
FROM
  aka,
  cast,
  company_name,
  movie_companies,
  name,
  role,
  title
WHERE
  company_name.country = 'usa' AND
  role.type = 'writer' AND
  aka.person_id = name.id AND
  cast.person_id = name.id AND
  aka.person_id = cast.person_id AND
  cast.movie_id = title.id AND
  movie_companies.movie = title.id_id AND
  cast.movie_id = movie_companies.movie_id AND
  movie_companies.company_id = company_name.id AND
  cast.role_id = role.id;

```

JOB Example Query



```

SELECT
    *
FROM
    aka,
    cast,
    company_name,
    movie_companies
WHERE
    aka.person_id = cast.person_id AND
    cast.movie_id = movie_companies.movie_id AND
    movie_companies.company_id = company_name.id;

```

Figure: Join Graph.

JOB Example Query




Figure: Join Graph.

```

SELECT
  *
FROM
  aka,
  cast,
  company_name,
  movie_companies
WHERE
  aka.person_id = cast.person_id AND
  cast.movie_id = movie_companies.movie_id AND
  movie_companies.company_id = company_name.id;


```

$$Q(x, y, z, w) :-$$

aka(x, y),
 cast(y, z),
 movie_companies(z, w),
 company_name(w)


Worst Case Scenario

$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$




Worst Case Scenario

$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$



A Better Plan

$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$



- 1 Query Optimization
- 2 Motivating Example
- 3 Prior Work: Cardinality Bounds**
- 4 Tightened Cardinality Bounds
- 5 Optimizations
- 6 Evaluation
 - Bound Tightening
 - Runtime Improvement
- 7 Conclusion and Future Directions

Review: Entropy

Take random variable X :

$$h(X) = - \sum_a \mathbb{P}(X = a) \cdot \log(\mathbb{P}(X = a))$$

Multiple variables:

$$h(X, Y) = - \sum_{a,b} \mathbb{P}(X = a, Y = b) \cdot \log(\mathbb{P}(X = a, Y = b))$$

Conditional Entropy:

$$h(X|Y) = - \sum_{a,b} \mathbb{P}(X = a, Y = b) \cdot \log\left(\frac{\mathbb{P}(X = a, Y = b)}{\mathbb{P}(Y = b)}\right)$$

Review: Entropy

Let X be uniformly distributed on the space $\{a_1, a_2, \dots, a_n\}$.

$$\begin{aligned}h(X) &= - \sum_{i=1}^n \mathbb{P}(X = a_i) \cdot \log(\mathbb{P}(X = a_i)) \\&= - \sum_{i=1}^n \frac{1}{n} \cdot \log\left(\frac{1}{n}\right) \\&= -n \frac{1}{n} \cdot \log\left(\frac{1}{n}\right) \\&= \log(n)\end{aligned}$$

Connection to Entropy

$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$

Connection to Entropy

$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$

- ▶ Create a random variable for each of the attributes present in the query.

$$x \rightarrow X, \quad y \rightarrow Y, \quad z \rightarrow Z, \quad w \rightarrow W$$

Connection to Entropy

$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$

- ▶ Create a random variable for each of the attributes present in the query.

$$x \rightarrow X, \quad y \rightarrow Y, \quad z \rightarrow Z, \quad w \rightarrow W$$

- ▶ Let (X, Y, Z, W) be uniformly distributed over all tuples in the true output of Q .

$$|Q(x, y, z, w)| = \exp(h(X, Y, Z, W))$$

Entropic Bounds

$$|Q(x, y, z, w)| = \exp(h(X, Y, Z, W))$$

- Suffices to bound $h(X, Y, Z, W)$.

Entropic Bounds

$$|Q(x, y, z, w)| = \exp(h(X, Y, Z, W))$$

- ▶ Suffices to bound $h(X, Y, Z, W)$.
- ▶ There are plenty of entropic bounds to choose from!

Entropic Bounds

$$h(X, Y, Z, W) \leq \dots$$

- 1 $h(X, Y) + h(Z|Y) + h(W|Z)$
 - 2 $h(X, Y) + h(Z|Y) + h(W)$
 - 3 $h(X, Y) + h(Z, W)$
 - 4 $h(X, Y) + h(Z|W) + h(W)$
 - 5 $h(X|Y) + h(Y, Z) + h(W|Z)$
 - 6 $h(X|Y) + h(Y, Z) + h(W)$
 - 7 $h(X|Y) + h(Y|Z) + h(Z, W)$
 - 8 $h(X|Y) + h(Y|Z) + h(Z|W) + h(Z)$
-

(only a subset of all entropic bounding formulas)

Entropic Bounds

$$h(X, Y, Z, W) \leq \dots$$

- 1 $h(X, Y) + h(Z|Y) + h(W|Z)$
 - 2 $h(X, Y) + h(Z|Y) + h(W)$
 - 3 $h(X, Y) + h(Z, W)$
 - 4 $h(X, Y) + h(Z|W) + h(W)$
 - 5 $h(X|Y) + h(Y, Z) + h(W|Z)$
 - 6 $h(X|Y) + h(Y, Z) + h(W)$
 - 7 $h(X|Y) + h(Y|Z) + h(Z, W)$
 - 8 $h(X|Y) + h(Y|Z) + h(Z|W) + h(Z)$
-

(only a subset of all entropic bounding formulas)

Entropic Bounds

$$|Q(x, y, z, w)| = \exp(h(X, Y, Z, W)) \leq \exp(h(X|Y) + h(Y, Z) + h(W|Z))$$

$$h(X|Y) \leq \log d_{\text{aka}}^y$$

$$h(Y, Z) \leq \log c_{\text{cast}}$$

$$h(W|Z) \leq \log d_{\text{movie_companies}}^z$$

$$d_{\text{aka}}^y = \text{“Max Degree”}$$

= Count of most common y attribute value in aka_name.

$$c_{\text{cast}} = \text{“Count”}$$

= Count of entire cast_info relation.

Cardinality Bound

$$\begin{aligned}
 |Q(x, y, z, w)| &= \exp(h(X, Y, Z, W)) \\
 &\leq \exp(\underbrace{h(X|Y)}_{\leq \log d_{\text{aka}}^y} + \underbrace{h(Y, Z)}_{\leq \log c_{\text{cast}}} + \underbrace{h(W|Z)}_{\leq \log d_{\text{movie_companies}}^z}) \\
 &\leq d_{\text{aka}}^y \cdot c_{\text{cast}} \cdot d_{\text{movie_companies}}^z
 \end{aligned}$$

Bound Formula Generation

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), \underbrace{movie_companies(z, w)}_{mc}, \underbrace{company_name(w)}_{cn}$$

x	y	z	w	entropic formula	bound formula
aka	aka	cast	mc	$h(X, Y) + h(Z Y) + h(W Z)$	$c_{aka} \cdot d_{cast}^y \cdot d_{mc}^z$
aka	aka	cast	cn	$h(X, Y) + h(Z Y) + h(W)$	$c_{aka} \cdot d_{cast}^y \cdot c_{cn}$
aka	aka	mc	mc	$h(X, Y) + h(Z, W)$	$c_{aka} \cdot c_{mc}$
aka	aka	mc	cn	$h(X, Y) + h(Z W) + h(W)$	$c_{aka} \cdot d_{mc}^w \cdot c_{cn}$
aka	cast	cast	mc	$h(X Y) + h(Y, Z) + h(W Z)$	$d_{aka}^y \cdot c_{cast} \cdot d_{mc}^z$
aka	cast	cast	cn	$h(X Y) + h(Y, Z) + h(W)$	$d_{aka}^y \cdot c_{cast} \cdot c_{cn}$
aka	cast	mc	mc	$h(X Y) + h(Y Z) + h(Z, W)$	$d_{aka}^y \cdot d_{cast}^z \cdot c_{mc}$
aka	cast	mc	cn	$h(X Y) + h(Y Z) + h(Z W) + h(Z)$	$d_{aka}^y \cdot d_{cast}^z \cdot d_{mc}^w \cdot c_{cn}$

Bound Formula Generation

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), \underbrace{movie_companies(z, w)}_{mc}, \underbrace{company_name(w)}_{cn}$$

x	y	z	w	entropic formula	bound formula
aka	aka	cast	mc	$h(X, Y) + h(Z Y) + h(W Z)$	$c_{aka} \cdot d_{cast}^y \cdot d_{mc}^z$
aka	aka	cast	cn	$h(X, Y) + h(Z Y) + h(W)$	$c_{aka} \cdot d_{cast}^y \cdot c_{cn}$
aka	aka	mc	mc	$h(X, Y) + h(Z, W)$	$c_{aka} \cdot c_{mc}$
aka	aka	mc	cn	$h(X, Y) + h(Z W) + h(W)$	$c_{aka} \cdot d_{mc}^w \cdot c_{cn}$
aka	cast	cast	mc	$h(X Y) + h(Y, Z) + h(W Z)$	$d_{aka}^y \cdot c_{cast} \cdot d_{mc}^z$
aka	cast	cast	cn	$h(X Y) + h(Y, Z) + h(W)$	$d_{aka}^y \cdot c_{cast} \cdot c_{cn}$
aka	cast	mc	mc	$h(X Y) + h(Y Z) + h(Z, W)$	$d_{aka}^y \cdot d_{cast}^z \cdot c_{mc}$
aka	cast	mc	cn	$h(X Y) + h(Y Z) + h(Z W) + h(Z)$	$d_{aka}^y \cdot d_{cast}^z \cdot d_{mc}^w \cdot c_{cn}$

Bound Formula Generation

$$Q(x, y, z, w) :- \text{aka}(x, y), \text{cast}(y, z), \underbrace{\text{movie_companies}(z, w)}_{\text{mc}}, \underbrace{\text{company_name}(w)}_{\text{cn}}$$

x	y	z	w	entropic formula	bound formula
aka	aka	cast	mc	$h(X, Y) + h(Z Y) + h(W Z)$	$C_{\text{aka}} \cdot d_{\text{cast}}^y \cdot d_{\text{mc}}^z$
aka	aka	cast	cn	$h(X, Y) + h(Z Y) + h(W)$	$C_{\text{aka}} \cdot d_{\text{cast}}^y \cdot C_{\text{cn}}$
aka	aka	mc	mc	$h(X, Y) + h(Z, W)$	$C_{\text{aka}} \cdot C_{\text{mc}}$
aka	aka	mc	cn	$h(X, Y) + h(Z W) + h(W)$	$C_{\text{aka}} \cdot d_{\text{mc}}^w \cdot C_{\text{cn}}$
aka	cast	cast	mc	$h(X Y) + h(Y, Z) + h(W Z)$	$d_{\text{aka}}^y \cdot C_{\text{cast}} \cdot d_{\text{mc}}^z$
aka	cast	cast	cn	$h(X Y) + h(Y, Z) + h(W)$	$d_{\text{aka}}^y \cdot C_{\text{cast}} \cdot C_{\text{cn}}$
aka	cast	mc	mc	$h(X Y) + h(Y Z) + h(Z, W)$	$d_{\text{aka}}^y \cdot d_{\text{cast}}^z \cdot C_{\text{mc}}$
aka	cast	mc	cn	$h(X Y) + h(Y Z) + h(Z W) + h(Z)$	$d_{\text{aka}}^y \cdot d_{\text{cast}}^z \cdot d_{\text{mc}}^w \cdot C_{\text{cn}}$

Bound Formula Generation

$$Q(x, y, z, w) :- \text{aka}(x, y), \text{cast}(y, z), \underbrace{\text{movie_companies}(z, w)}_{\text{mc}}, \underbrace{\text{company_name}(w)}_{\text{cn}}$$

x	y	z	w	entropic formula	bound formula
aka	aka	cast	mc	$h(X, Y) + h(Z Y) + h(W Z)$	$c_{\text{aka}} \cdot d_{\text{cast}}^y \cdot d_{\text{mc}}^z$
aka	aka	cast	cn	$h(X, Y) + h(Z Y) + h(W)$	$c_{\text{aka}} \cdot d_{\text{cast}}^y \cdot c_{\text{cn}}$
aka	aka	mc	mc	$h(X, Y) + h(Z, W)$	$c_{\text{aka}} \cdot c_{\text{mc}}$
aka	aka	mc	cn	$h(X, Y) + h(Z W) + h(W)$	$c_{\text{aka}} \cdot d_{\text{mc}}^w \cdot c_{\text{cn}}$
aka	cast	cast	mc	$h(X Y) + h(Y, Z) + h(W Z)$	$d_{\text{aka}}^y \cdot c_{\text{cast}} \cdot d_{\text{mc}}^z$
aka	cast	cast	cn	$h(X Y) + h(Y, Z) + h(W)$	$d_{\text{aka}}^y \cdot c_{\text{cast}} \cdot c_{\text{cn}}$
aka	cast	mc	mc	$h(X Y) + h(Y Z) + h(Z, W)$	$d_{\text{aka}}^y \cdot d_{\text{cast}}^z \cdot c_{\text{mc}}$
aka	cast	mc	cn	$h(X Y) + h(Y Z) + h(Z W) + h(Z)$	$d_{\text{aka}}^y \cdot d_{\text{cast}}^z \cdot d_{\text{mc}}^w \cdot c_{\text{cn}}$

Entropic Bounds

Neat! But is it useful?

- ▶ Short answer: No. (Not yet, anyway)

Entropic Bounds

Neat! But is it useful?

- ▶ Short answer: No. (Not yet, anyway)
 - ▶ Bounds are still far too loose (overestimation).

Entropic Bounds

Neat! But is it useful?

- ▶ Short answer: No. (Not yet, anyway)
 - ▶ Bounds are still far too loose (overestimation).
 - ▶ Need to tighten the bounds.

Entropic Bounds

Neat! But is it useful?

- ▶ Short answer: No. (Not yet, anyway)
 - ▶ Bounds are still far too loose (overestimation).
 - ▶ Need to tighten the bounds.
- ▶ How to tighten? Partitioning.

- 1 Query Optimization
- 2 Motivating Example
- 3 Prior Work: Cardinality Bounds
- 4 Tightened Cardinality Bounds**
- 5 Optimizations
- 6 Evaluation
 - Bound Tightening
 - Runtime Improvement
- 7 Conclusion and Future Directions

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$

aka

cast

movie_companies


company_name

aka

cast

mc

cn

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$


Hash the values of each tuple and bucketize on the hash values.

$$aka[1, 0] = \{t \in cast \mid hash(t[y]) = 1 \wedge hash(t[z]) = 0\}$$

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$

aka

cast

movie_companies

company_name

cast[1,0]

aka[0,1]


mc[0,1]

cn[1]


- Pick a hash value for each attribute in the query:

$$x, y, z, w \rightarrow [0, 1, 0, 1]$$

- The matching buckets from each relation is the partition $D[0, 1, 0, 1]$.

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$


- ▶ $Q(D)$: query evaluated on database D .
- ▶ $Q(D[J])$: query evaluated on partition $D[J]$.

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$


- ▶ Bound each partition $D[J]$.
- ▶ Sum will be a bound on the full database D .


$$Q(D) = \bigcup_J Q(D[J])$$

$$|Q(D)| \leq \sum_J bound(Q(D[J]))$$

Partition Bounding


$$|Q(D)| \leq \sum_{J \in \{0,1\}^4} \min \left\{ \begin{array}{l} c_{aka}[J] \cdot d_{cast}^y[J] \cdot d_{mc}^z[J] \\ c_{aka}[J] \cdot d_{cast}^y[J] \cdot c_{cn}[J] \\ c_{aka}[J] \cdot c_{mc}[J] \\ c_{aka}[J] \cdot d_{mc}^w[J] \cdot c_{cn}[J] \\ d_{aka}^y[J] \cdot c_{cast}[J] \cdot d_{mc}^z[J] \\ d_{aka}^y[J] \cdot c_{cast}[J] \cdot c_{cn}[J] \\ d_{aka}^y[J] \cdot d_{cast}^z[J] \cdot c_{mc}[J] \\ d_{aka}^y[J] \cdot d_{cast}^z[J] \cdot d_{mc}^w[J] \cdot c_{cn}[J] \end{array} \right.$$

The Bound Sketch




- One bound sketch per table.

The Bound Sketch




- ▶ One bound sketch per table.
- ▶ Need count and degree statistics.

The Bound Sketch




- ▶ One bound sketch per table.
- ▶ Need count and degree statistics.
- ▶ Some calculated offline, some at runtime.

The Bound Sketch



- ▶ One bound sketch per table.
- ▶ Need count and degree statistics.
- ▶ Some calculated offline, some at runtime.
- ▶ Like a richer randomized histogram.

The Bound Sketch



- ▶ One bound sketch per table.
- ▶ Need count and degree statistics.
- ▶ Some calculated offline, some at runtime.
- ▶ Like a richer randomized histogram.
- ▶ Restriction: this method is only suitable for equijoins.

Exponential Growth

- ▶ Sketch size (number of buckets) exponential in hash size.
 - ▶ Exponent is number of attributes in relation.

Exponential Growth

- ▶ Sketch size (number of buckets) exponential in hash size.
 - ▶ Exponent is number of attributes in relation.
- ▶ Number of elements to sum up exponential in hash size.
 - ▶ Exponent is number of attributes in entire query.

Exponential Growth

- ▶ Sketch size (number of buckets) exponential in hash size.
 - ▶ Exponent is number of attributes in relation.
- ▶ Number of elements to sum up exponential in hash size.
 - ▶ Exponent is number of attributes in entire query.
- ▶ Non-monotonic bound behavior

Tuning Bucket Allocation

- ▶ As the number of buckets increases, we get more information, and bounds tighten, right?

Tuning Bucket Allocation

- ▶ As the number of buckets increases, we get more information, and bounds tighten, right?
- ▶ When exclusively partitioning unconditionally covered attributes: yes.

Tuning Bucket Allocation

- ▶ As the number of buckets increases, we get more information, and bounds tighten, right?
- ▶ When exclusively partitioning unconditionally covered attributes: yes.
- ▶ When also partitioning conditionally covered attributes: not necessarily.

Tuning Bucket Allocation

- ▶ As the number of buckets increases, we get more information, and bounds tighten, right?
- ▶ When exclusively partitioning unconditionally covered attributes: yes.
- ▶ When also partitioning conditionally covered attributes: not necessarily.
 - ▶ Non-monotonic tradeoff space.

Example of Non-monotonic Behavior

$$Q(x, y, z) :- R(z, y), S(y, z), T(z, w)$$

x	y		y	z		z	w		x	y	z	w
0	0		0	0		0	0		0	0	0	0
0	1	\bowtie	1	0	\bowtie	1	1	$=$	0	1	0	0
1	0		2	1		2	2		1	0	0	0
1	1		3	1		3	3		1	1	0	0
R			S			T			Q			

Example of Non-monotonic Behavior

x	y		y	z		z	w		x	y	z	w
0	0		0	0		0	0		0	0	0	0
0	1	⋈	1	0	⋈	1	1	=	0	1	0	0
1	0		2	1		2	2		1	0	0	0
1	1		3	1		3	3		1	1	0	0
R			S			T			Q			

Example of Non-monotonic Behavior

x	y		y	z		z	w		x	y	z	w
0	0		0	0		0	0		0	0	0	0
0	1	⋈	1	0	⋈	1	1	=	0	1	0	0
1	0		2	1		2	2		1	0	0	0
1	1		3	1		3	3		1	1	0	0
R			S			T			Q			

$$|Q(x, y, z)| \leq \min \begin{cases} c_R \cdot d_S^y \cdot d_T^z \\ d_R^y \cdot c_S \cdot d_T^z \\ d_R^y \cdot d_S^z \cdot c_T \\ c_R \cdot c_T \end{cases}$$

Example of Non-monotonic Behavior

x	y		y	z		z	w		x	y	z	w
0	0		0	0		0	0		0	0	0	0
0	1	⋈	1	0	⋈	1	1	=	0	1	0	0
1	0		2	1		2	2		1	0	0	0
1	1		3	1		3	3		1	1	0	0
R			S			T			Q			

$$|Q(x, y, z)| \leq \min \begin{cases} c_R \cdot d_S^y \cdot d_T^z \\ d_R^y \cdot c_S \cdot d_T^z \\ d_R^y \cdot d_S^z \cdot c_T \\ c_R \cdot c_T \end{cases}$$

$$c_{R(0)} \cdot d_{S(0,0)}^y \cdot d_{T(0)}^z = 4 \cdot 1 \cdot 1 = 4$$

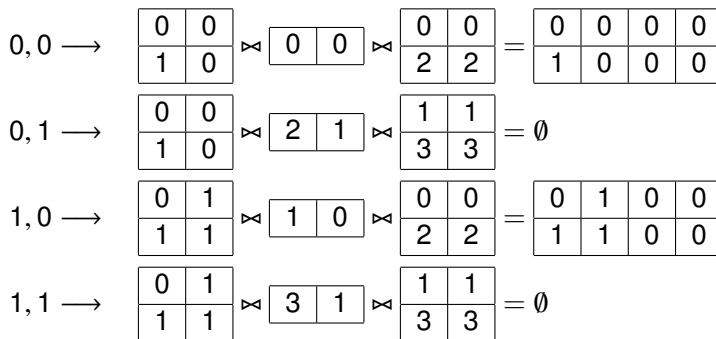
- ▶ Define hash function $\text{hash}(u_i) = i \% 2$.

$$\text{hash}(0) = \text{hash}(2) = 0$$

$$\text{hash}(1) = \text{hash}(3) = 1$$

Partitioned Relations

$\text{hash}(y), \text{hash}(z) = \dots$



$$\begin{array}{lcl}
0,0 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 2 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 0 & 0 & 0 \\ \hline 1 & 0 & 0 & 0 \\ \hline \end{array} \\
0,1 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 2 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 1 \\ \hline 3 & 3 \\ \hline \end{array} = \emptyset \\
1,0 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 1 \\ \hline 1 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 2 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 1 & 0 & 0 \\ \hline 1 & 1 & 0 & 0 \\ \hline \end{array} \\
1,1 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 1 \\ \hline 1 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 3 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 1 \\ \hline 3 & 3 \\ \hline \end{array} = \emptyset
\end{array}$$

$$\sum_{\substack{i,j \\ \in \{0,1\}}} \min \begin{cases} c_{R(i)} \cdot d_{S(i,j)}^y \cdot d_{T(j)}^z \\ d_{R(i)}^y \cdot c_{S(i,j)} \cdot d_{T(j)}^z \\ d_{R(i)}^y \cdot d_{S(i,j)}^z \cdot c_{T(j)} \\ c_{R(i)} \cdot c_{T(j)} \end{cases}$$

$$\begin{array}{lcl}
0,0 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 2 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 0 & 0 & 0 \\ \hline 1 & 0 & 0 & 0 \\ \hline \end{array} \\
0,1 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 2 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 1 \\ \hline 3 & 3 \\ \hline \end{array} = \emptyset \\
1,0 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 1 \\ \hline 1 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 2 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 1 & 0 & 0 \\ \hline 1 & 1 & 0 & 0 \\ \hline \end{array} \\
1,1 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 1 \\ \hline 1 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 3 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 1 \\ \hline 3 & 3 \\ \hline \end{array} = \emptyset
\end{array}$$

$$\sum_{\substack{i,j \\ \in \{0,1\}}} \min \begin{cases} c_{R(i)} \cdot d_{S(i,j)}^y \cdot d_{T(j)}^z \\ d_{R(i)}^y \cdot c_{S(i,j)} \cdot d_{T(j)}^z \\ d_{R(i)}^y \cdot d_{S(i,j)}^z \cdot c_{T(j)} \\ c_{R(i)} \cdot c_{T(j)} \end{cases} = \sum_{\substack{i,j \\ \in \{0,1\}}} \min \begin{cases} 2 \cdot 1 \cdot 1 \\ 2 \cdot 1 \cdot 1 \\ 2 \cdot 1 \cdot 2 \\ 2 \cdot 2 \end{cases}$$

$$\begin{array}{lcl}
0,0 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 2 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 0 & 0 & 0 \\ \hline 1 & 0 & 0 & 0 \\ \hline \end{array} \\
0,1 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 2 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 1 \\ \hline 3 & 3 \\ \hline \end{array} = \emptyset \\
1,0 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 1 \\ \hline 1 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 0 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 2 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 1 & 0 & 0 \\ \hline 1 & 1 & 0 & 0 \\ \hline \end{array} \\
1,1 \longrightarrow & \begin{array}{|c|c|} \hline 0 & 1 \\ \hline 1 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 3 & 1 \\ \hline \end{array} \bowtie \begin{array}{|c|c|} \hline 1 & 1 \\ \hline 3 & 3 \\ \hline \end{array} = \emptyset
\end{array}$$

$$\sum_{\substack{i,j \\ \in \{0,1\}}} \min \begin{cases} c_{R(i)} \cdot d_{S(i,j)}^y \cdot d_{T(j)}^z \\ d_{R(i)}^y \cdot c_{S(i,j)} \cdot d_{T(j)}^z \\ d_{R(i)}^y \cdot d_{S(i,j)}^z \cdot c_{T(j)} \\ c_{R(i)} \cdot c_{T(j)} \end{cases} = \sum_{\substack{i,j \\ \in \{0,1\}}} \min \begin{cases} 2 \cdot 1 \cdot 1 \\ 2 \cdot 1 \cdot 1 \\ 2 \cdot 1 \cdot 2 \\ 2 \cdot 2 \end{cases} = \sum_{\substack{i,j \\ \in \{0,1\}}} 2 = 8$$

Non-Linearity of Degree Statistic

- ▶ Count is linear with respect to disjoint union!

$$\text{count}(A) + \text{count}(B) = \text{count}(A \cup B)$$

- ▶ Degree is not...

$$\text{degree}(A) + \text{degree}(B) \geq \text{degree}(A \cup B)$$

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$

aka

cast

movie_companies

company_name

aka

cast

mc

cn

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$

$$c_{aka} \cdot d_{cast}^y \cdot d_{mc}^z$$

aka

cast

movie_companies

company_name

aka

cast

mc

cn

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$



$$c_{aka} \cdot d_{cast}^y \cdot d_{mc}^z$$

aka

cast

movie_companies

company_name



$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$

$$d_{aka}^y \cdot c_{cast} \cdot d_{mc}^z$$

aka

cast

movie_companies

company_name

aka

cast

mc

cn

$$Q(x, y, z, w) :- aka(x, y), cast(y, z), movie_companies(z, w), company_name(w)$$



$$d_{aka}^y \cdot c_{cast} \cdot d_{mc}^z$$

aka

cast

movie_companies

company_name



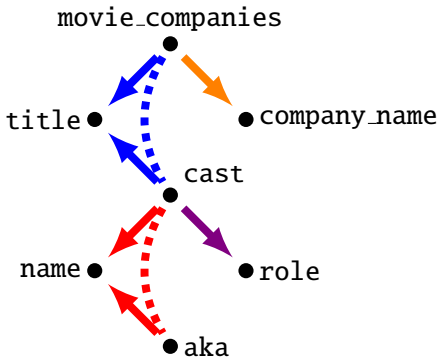
Bound Calculation

Calculate the minimal bound over all entropic bounds.

$$|Q(D)| \leq \min_{\substack{b \in \\ \text{bounding formulas}}} \left(\sum_{\substack{J \in \\ \text{partition indexes}}} b(Q(D[J])) \right)$$

- 1 Query Optimization
- 2 Motivating Example
- 3 Prior Work: Cardinality Bounds
- 4 Tightened Cardinality Bounds
- 5 Optimizations**
- 6 Evaluation
 - Bound Tightening
 - Runtime Improvement
- 7 Conclusion and Future Directions


Filter Predicate Analysis



```

SELECT
    *
FROM
    aka,
    cast,
    company_name,
    movie_companies,
    name,
    role,
    title
WHERE
    company_name.country = 'usa' AND
    role.type = 'writer' AND
    aka.person_id = name.id AND
    cast.person_id = name.id AND
    aka.person_id = cast.person_id AND
    cast.movie_id = title.id AND
    movie_companies.movie = title.id_id AND
    cast.movie_id = movie_companies.movie_id AND
    movie_companies.company_id = company_name.id AND
    cast.role_id = role.id;
  
```


Filter Predicate Analysis



```

SELECT
    *
FROM
    aka,
    cast,
    company_name,
    movie_companies,
    name,
    role,
    title
WHERE
    company_name.country = 'usa' AND
    role.type = 'writer' AND
    aka.person_id = name.id AND
    cast.person_id = cast.person_id AND
    aka.person_id = cast.person_id AND
    cast.movie_id = title.id AND
    movie_companies.movie = title.id_id AND
    cast.movie_id = movie_companies.movie_id AND
    movie_companies.company_id = company_name.id AND
    cast.role_id = role.id;
  
```

Filter Predicate Analysis




```

SELECT
    *
FROM
    aka,
    cast,
    company_name,
    movie_companies,
    name,
    role,
    title
WHERE
    company_name.country = 'usa' AND
    role.type = 'writer' AND
    aka.person.id = name.id AND
    cast.person.id = cast.person.id AND
    cast.movie.id = title.id AND
    movie_companies.movie = title.id.id AND
    cast.movie.id = movie_companies.movie.id AND
    movie_companies.company.id = company_name.id AND
    cast.role.id = role.id;

```

Filter Predicate Analysis



```

SELECT
    *
FROM
    aka,
    cast,
    company_name,
    movie_companies,
    name,
    role,
    title
WHERE
    company_name.country = 'usa' AND
    role.type = 'writer' AND
    aka.person.id = name.id AND
    cast.person.id = cast.person.id AND
    cast.movie.id = title.id AND
    movie_companies.movie = title.id.id AND
    cast.movie.id = movie_companies.movie.id AND
    movie_companies.company.id = company_name.id AND
    cast.role.id = role.id;

```


Filter Propagation




Figure: Original hypergraph representation.

Filter Propagation




Figure: Original hypergraph representation.

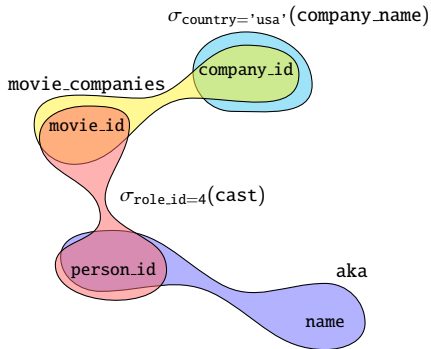


Figure: Hypergraph after selection propagation and elimination.

Table Scans During Optimization

Analysis of selection predicates can lead to:

Table Scans During Optimization

Analysis of selection predicates can lead to:

- ▶ Full propagation.
 - ▶ Highly selective predicate: yields fewer tuples than the hash size.
 - ▶ Scans on predicate relation and (most likely) on foreign key relation.

Table Scans During Optimization

Analysis of selection predicates can lead to:

- ▶ Full propagation.
 - ▶ Highly selective predicate: yields fewer tuples than the hash size.
 - ▶ Scans on predicate relation and (most likely) on foreign key relation.
- ▶ Updated the bound sketch.
 - ▶ Selective predicate but more tuples than hash size.
 - ▶ Scan on predicate relation.

Table Scans During Optimization

Analysis of selection predicates can lead to:

- ▶ Full propagation.
 - ▶ Highly selective predicate: yields fewer tuples than the hash size.
 - ▶ Scans on predicate relation and (most likely) on foreign key relation.
- ▶ Updated the bound sketch.
 - ▶ Selective predicate but more tuples than hash size.
 - ▶ Scan on predicate relation.
- ▶ Defaulting to unmodified bound sketch.
 - ▶ Non selective predicate.
 - ▶ Early exit during scan on predicate relation.

Table Scans During Optimization

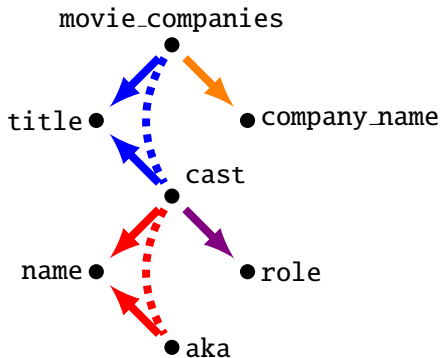



Table Scans During Optimization



- 1 Query Optimization
- 2 Motivating Example
- 3 Prior Work: Cardinality Bounds
- 4 Tightened Cardinality Bounds
- 5 Optimizations
- 6 Evaluation**
 - **Bound Tightening**
 - Runtime Improvement
- 7 Conclusion and Future Directions



Googleplus Microbenchmark Examples

```
SELECT COUNT(*)
FROM
  community_44 AS t0,
  community_44 AS t1,
  community_44 AS t2,
  community_44 AS t3
WHERE
  t0.object = t1.subject AND
  t1.object = t2.subject AND
  t2.object = t3.subject AND
  t0.subject % 512 = 89 AND
  t3.object % 512 = 174;
```

Googleplus Microbenchmark Examples

```
SELECT COUNT(*)
FROM
  community_30 AS t0,
  community_30 AS t1,
  community_30 AS t2,
  community_30 AS t3,
  community_30 AS t4
WHERE
  t0.object = t1.subject AND
  t0.object = t2.subject AND
  t0.object = t3.subject AND
  t3.object = t4.subject AND
  t0.subject % 256 = 49 AND
  t1.object % 256 = 213 AND
  t2.object % 256 = 152 AND
  t4.object % 256 = 248;
AND ci.movie_id = mc.movie_id;
```

Googleplus Progressive Bound Tightness




- 1 Query Optimization
- 2 Motivating Example
- 3 Prior Work: Cardinality Bounds
- 4 Tightened Cardinality Bounds
- 5 Optimizations
- 6 Evaluation**
 - Bound Tightening
 - **Runtime Improvement**
- 7 Conclusion and Future Directions


Join Order Benchmark

- ▶ Built on the IMDb dataset.
 - ▶ 113 queries.
 - ▶ 33 unique topologies.
 - ▶ Skew!
 - ▶ Correlation!
 - ▶ Complex selection predicates!
- ▶ *How Good Are Query Optimizers, Really?* Leis et al. VLDB 2015.


Bound Relative Error Versus Postgres Relative Error




Bound Q-Error Versus Postgres Q-Error



Plan Execution Runtime (With Foreign Keys Indexes)



Plan Execution Runtime (With Foreign Keys Indexes)



Plan Execution Runtime (No Foreign Key Indexes)



Plan Execution Runtime (No Foreign Key Indexes)



Figure: Linear scale plan execution time over JOB queries.

- ▶ Total runtime (including 1 hour cutoff for postgres).
 - ▶ Postgres: 21,125 seconds.
 - ▶ Bound (4096 buckets): 2,216 seconds.

Takeaways

- ▶ Significant gain for very slow queries.

Takeaways

- ▶ Significant gain for very slow queries.
- ▶ On par with fast queries.

- 1 Query Optimization
- 2 Motivating Example
- 3 Prior Work: Cardinality Bounds
- 4 Tightened Cardinality Bounds
- 5 Optimizations
- 6 Evaluation
 - Bound Tightening
 - Runtime Improvement
- 7 Conclusion and Future Directions**

Optimization Time

- ▶ Currently using naive enumeration and sketch construction approach.

Optimization Time

- ▶ Currently using naive enumeration and sketch construction approach.
- ▶ Approximation of degree statistics.

Contributions

- ▶ Technique for tightening theoretically guaranteed join cardinality upper bounds.

Contributions

- ▶ Technique for tightening theoretically guaranteed join cardinality upper bounds.
- ▶ Method for enumerating practical subset of bounding formulas.

Contributions

- ▶ Technique for tightening theoretically guaranteed join cardinality upper bounds.
- ▶ Method for enumerating practical subset of bounding formulas.
- ▶ Partition budgeting strategy to control the space complexity of our sketches, and the time complexity of our bound calculation.

Contributions

- ▶ Technique for tightening theoretically guaranteed join cardinality upper bounds.
- ▶ Method for enumerating practical subset of bounding formulas.
- ▶ Partition budgeting strategy to control the space complexity of our sketches, and the time complexity of our bound calculation.
- ▶ Demonstrate practicality on challenging real world benchmark.

Acknowledgements

- ▶ Thank you to Jenny, Tomer, Laurel, Brandon, Jingjing, Tobin, and Leilani!
- ▶ This research is supported by NSF grant AITF 1535565 and III 1614738.